

Assignment on Ensemble Learning



Course name: Data Science

Course code: CSEL – 42---

Submitted By

Farhana Akter Suci

ID: B190305001

&

Rifah Sajida Deya

ID: B190305004

Submitted To

Dr. Md. Manowarul Islam

Associate Professor, Department of C S E, Jagannath University

5th February , 2025

✓ 1) Load and Preprocess the Data:

```
import pandas as pd
```

```
data = pd.read_csv("Maternal Health Risk Data Set.csv")
```

data



	Age	SystolicBP	DiastolicBP	BS	BodyTemp	HeartRate	RiskLevel
0	25	130	80	15.0	98.0	86	high risk
1	35	140	90	13.0	98.0	70	high risk
2	29	90	70	8.0	100.0	80	high risk
3	30	140	85	7.0	98.0	70	high risk
4	35	120	60	6.1	98.0	76	low risk
...
1009	22	120	60	15.0	98.0	80	high risk
1010	55	120	90	18.0	98.0	60	high risk
1011	35	85	60	19.0	98.0	86	high risk
1012	43	120	90	18.0	98.0	70	high risk
1013	32	120	65	6.0	101.0	76	mid risk

1014 rows × 7 columns

```
data.shape
```



(1014, 7)

```
data.isnull().sum().sum()
```



0

✓ There is no null value in this dataset

```
data.info()
```

```

↳ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1014 entries, 0 to 1013
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              1014 non-null   int64
1   SystolicBP       1014 non-null   int64
2   DiastolicBP      1014 non-null   int64
3   BS               1014 non-null   float64
4   BodyTemp         1014 non-null   float64
5   HeartRate        1014 non-null   int64
6   RiskLevel        1014 non-null   object
dtypes: float64(2), int64(4), object(1)
memory usage: 55.6+ KB

```

```
data.describe().T
```

```

↳

```

	count	mean	std	min	25%	50%	75%	max
Age	1014.0	29.871795	13.474386	10.0	19.0	26.0	39.0	70.0
SystolicBP	1014.0	113.198225	18.403913	70.0	100.0	120.0	120.0	160.0
DiastolicBP	1014.0	76.460552	13.885796	49.0	65.0	80.0	90.0	100.0
BS	1014.0	8.725986	3.293532	6.0	6.9	7.5	8.0	19.0
BodyTemp	1014.0	98.665089	1.371384	98.0	98.0	98.0	98.0	103.0
HeartRate	1014.0	74.301775	8.088702	7.0	70.0	76.0	80.0	90.0

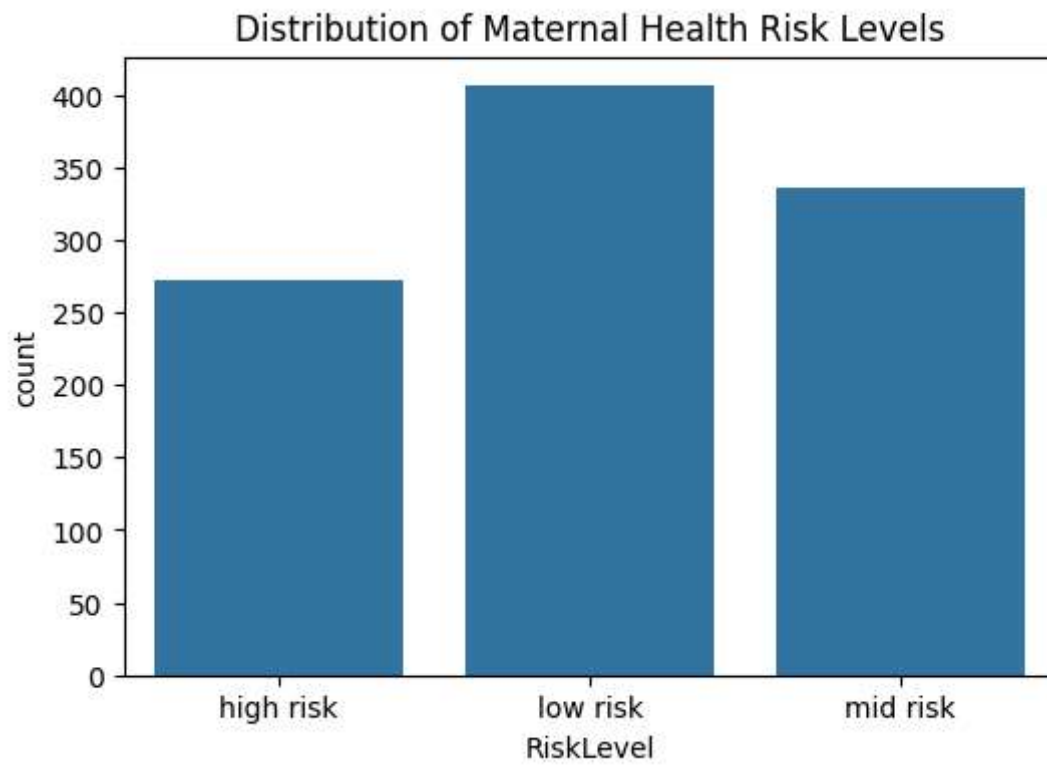
```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```

plt.figure(figsize=(6, 4))
sns.countplot(x="RiskLevel", data=data)
plt.title("Distribution of Maternal Health Risk Levels")
plt.show()

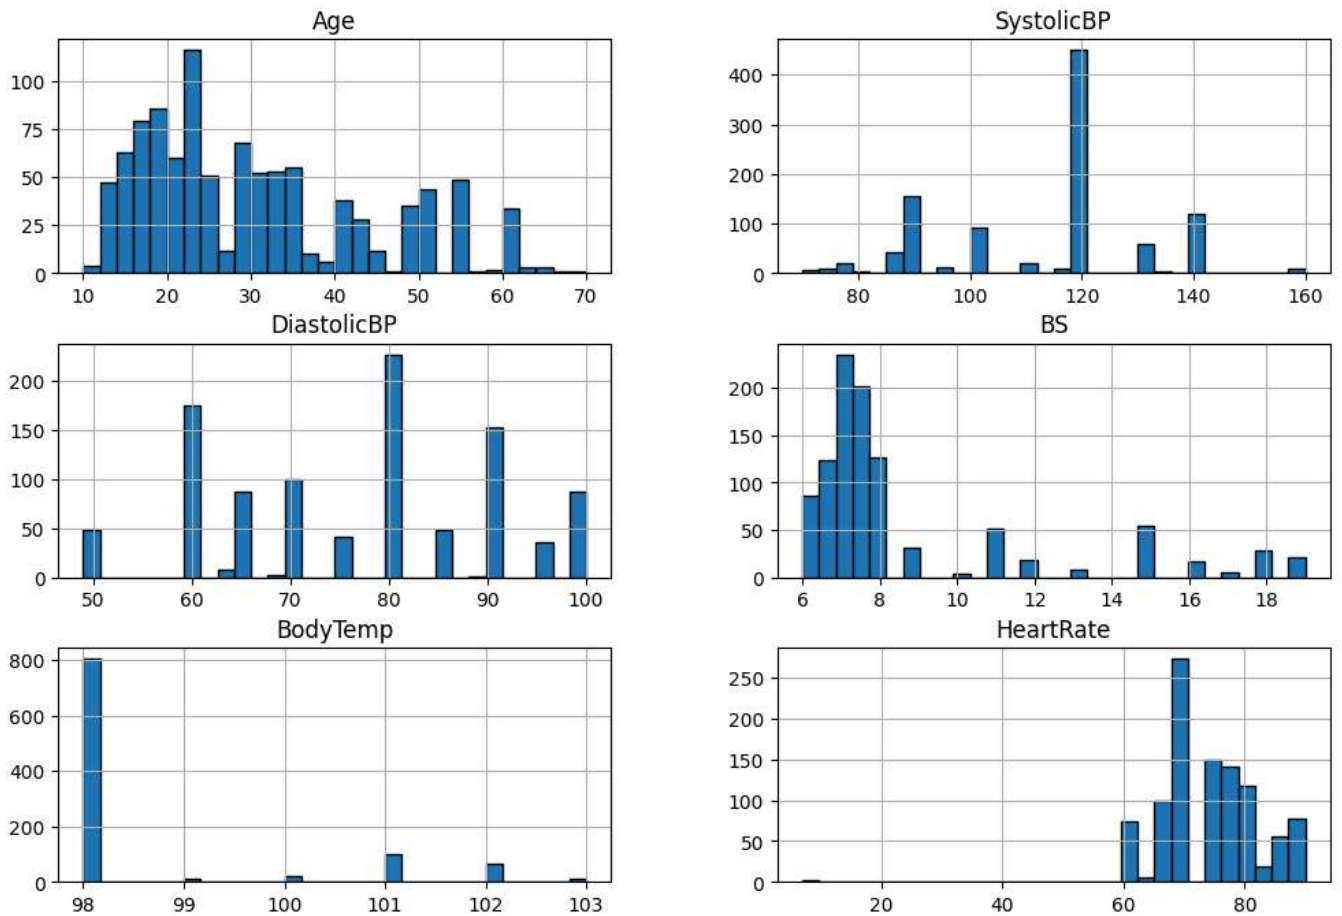
```



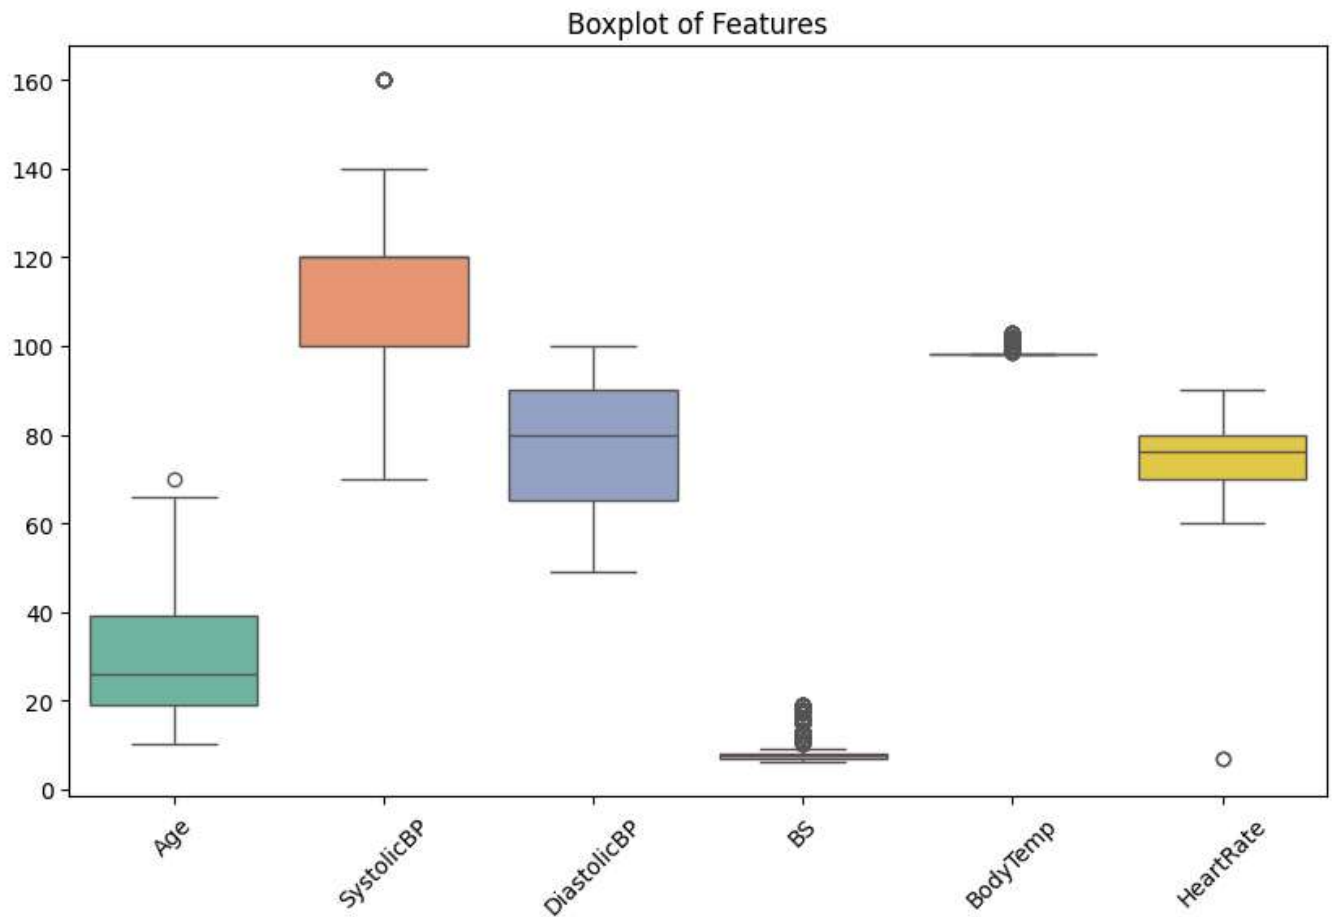
```
data.hist(figsize=(12, 8), bins=30, edgecolor='black')  
plt.suptitle("Feature Distributions", fontsize=14)  
plt.show()
```



Feature Distributions



```
plt.figure(figsize=(10, 6))
sns.boxplot(data=data, palette="Set2")
plt.xticks(rotation=45)
plt.title("Boxplot of Features")
plt.show()
```



✓ Label Encoding the target column


```
from sklearn.preprocessing import LabelEncoder
```

```
label_encoder=LabelEncoder()
```

```
data['RiskLevel']= label_encoder.fit_transform(data['RiskLevel'])
```

```
label_encoder = LabelEncoder()
```

```
plt.figure(figsize=(8, 6))  
sns.heatmap(data.corr(), annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)  
plt.title("Feature Correlation Heatmap")  
plt.show()
```



✓ Skewness Analysis

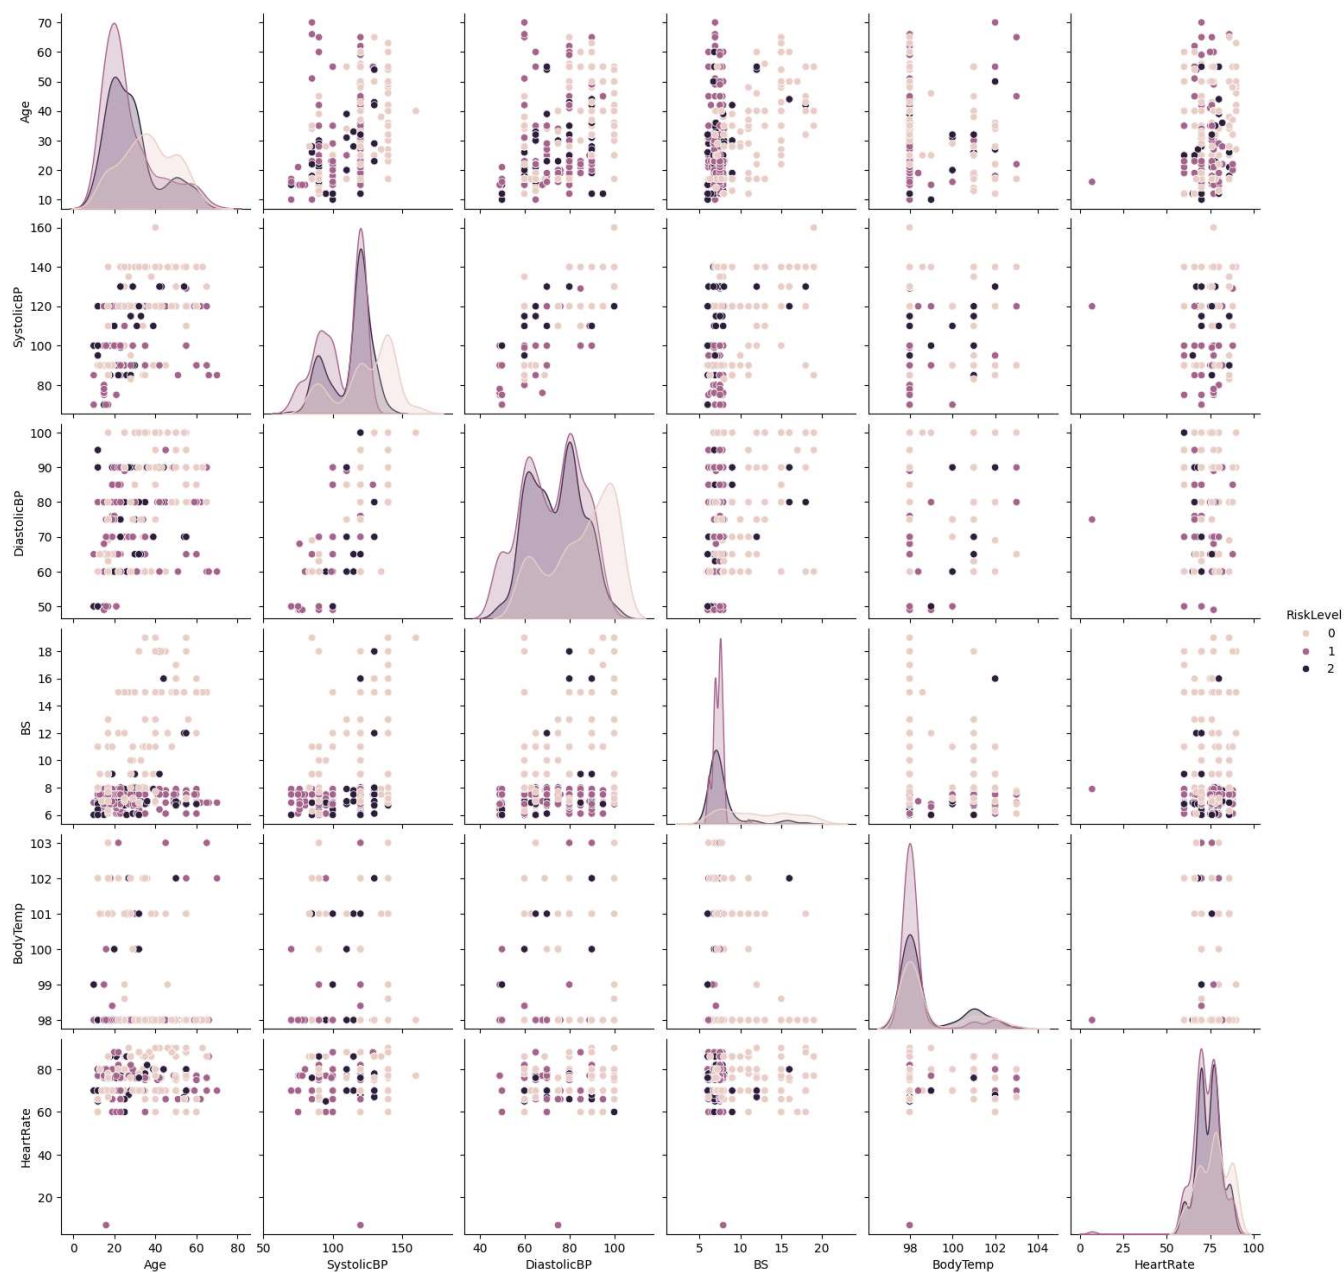
```
from scipy.stats import skew

skewness = data.skew()
print("Feature Skewness:\n", skewness)
```



```
Feature Skewness:
Age          0.783063
SystolicBP   -0.251189
DiastolicBP  -0.048441
BS           1.868203
BodyTemp     1.750988
HeartRate    -1.043525
RiskLevel    -0.108748
dtype: float64
```


```
sns.pairplot(data, hue="RiskLevel")  
plt.show()
```

✓ Split the dataset

```
X = data.iloc[:, :-1]
```

X



	Age	SystolicBP	DiastolicBP	BS	BodyTemp	HeartRate
0	25	130	80	15.0	98.0	86
1	35	140	90	13.0	98.0	70
2	29	90	70	8.0	100.0	80
3	30	140	85	7.0	98.0	70
4	35	120	60	6.1	98.0	76
...
1009	22	120	60	15.0	98.0	80
1010	55	120	90	18.0	98.0	60
1011	35	85	60	19.0	98.0	86
1012	43	120	90	18.0	98.0	70
1013	32	120	65	6.0	101.0	76

1014 rows × 6 columns

```
y = data.iloc[:, -1]
```

y

```
⇒ 0      0
   1      0
   2      0
   3      0
   4      1
   ..
 1009     0
 1010     0
 1011     0
 1012     0
 1013     2
    Name: RiskLevel, Length: 1014, dtype: int64
```

```
y.unique()
```

```
⇒ array([0, 1, 2])
```

```
y.nunique()
```

```
⇒ 3
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X = scaler.fit_transform(X)
```

✓ Label Encoding the target column

```
y = data['RiskLevel']
```

y

```
⇒ 0      0
   1      0
   2      0
   3      0
   4      1
   ..
 1009     0
```

```
1010    0
1011    0
1012    0
1013    2
Name: RiskLevel, Length: 1014, dtype: int64
```

✓ Split dataset: 80% training, 20% testing

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
→ ((811, 6), (203, 6), (811,), (203,))
```

Double-click (or enter) to edit

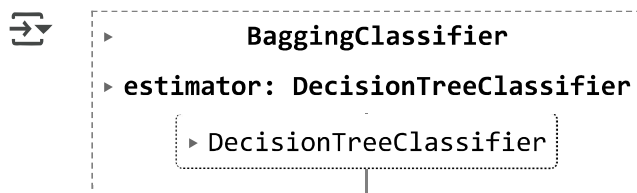
✓ 2) Bagging Implementation:

```
from sklearn.ensemble import BaggingClassifier
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
model1 = BaggingClassifier(estimator=DecisionTreeClassifier(), n_estimators=50, random_state
```

```
model1.fit(X_train, y_train)
```



```
test_predictions = model1.predict(X_test)
```

```
test_predictions
```

```
→ array([2, 0, 0, 1, 1, 2, 2, 2, 2, 1, 0, 2, 2, 0, 2, 2, 0, 1, 1, 2, 2, 2,
         0, 0, 1, 2, 0, 2, 1, 1, 2, 0, 1, 2, 1, 0, 0, 2, 1, 0, 0, 1, 0, 0,
         1, 2, 0, 2, 1, 1, 2, 1, 1, 2, 2, 2, 0, 1, 1, 2, 1, 2, 2, 2, 1, 0,
```

```
2, 0, 2, 0, 0, 0, 2, 2, 2, 0, 1, 2, 0, 2, 1, 2, 1, 1, 2, 1, 1, 2,
0, 2, 2, 0, 1, 0, 2, 1, 2, 1, 0, 2, 1, 2, 1, 1, 2, 2, 1, 2, 1, 2,
2, 2, 1, 0, 2, 1, 1, 0, 2, 0, 2, 1, 1, 2, 0, 0, 2, 2, 2, 2, 1, 1,
2, 2, 1, 1, 0, 2, 1, 2, 0, 2, 0, 0, 1, 1, 0, 1, 1, 2, 1, 2, 1, 2,
0, 2, 1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 2, 2, 1, 1, 0, 1, 2, 1, 2, 0,
2, 1, 1, 1, 0, 2, 1, 2, 0, 1, 0, 1, 2, 1, 2, 2, 2, 2, 1, 2, 2, 0,
1, 2, 2, 2, 0])
```

```
from sklearn.metrics import accuracy_score
```

```
from sklearn.metrics import precision_score, recall_score, f1_score, confusion_matrix
```

```
accuracy1=accuracy_score(y_test, test_predictions)
accuracy1
```

```
⇒ 0.8078817733990148
```

```
precision1=precision_score(y_test, test_predictions, average='weighted')
precision1
```

```
⇒ 0.814527289486854
```

```
recall1=recall_score(y_test, test_predictions, average='weighted')
recall1
```

```
⇒ 0.8078817733990148
```

```
f11=f1_score(y_test, test_predictions, average='weighted')
f11
```

```
⇒ 0.8084290104530727
```

✓ Show accuracy matrix as a Dataframe

```
accuracy_matrix = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score"],
    "Score": [accuracy1, precision1, recall1, f11]
})
```

```
accuracy_matrix
```



	Metric	Score
0	Accuracy	0.807882
1	Precision	0.814527
2	Recall	0.807882
3	F1-Score	0.808429

✓ Compute confusion matrix

```
conf_matrix1 = confusion_matrix(y_test, test_predictions)
```

```
conf_matrix1
```




```
array([[40,  1,  6],  
       [ 3, 60, 17],  
       [ 3,  9, 64]])
```

✓ Plot confusion matrix

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(6, 5))  
sns.heatmap(conf_matrix1, annot=True)  
plt.xlabel("Predicted Labels")  
plt.ylabel("True Labels")  
plt.title("Confusion Matrix")  
plt.show()
```



3) Boosting Implementation:

```
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier
```

```
from xgboost import XGBClassifier
```

```
model2=AdaBoostClassifier(n_estimators=50, random_state=42)
```

```
model3= GradientBoostingClassifier(n_estimators=50, random_state=42)
```

```
model4=XGBClassifier(n_estimators=50, eval_metric='mlogloss', random_state=42)
```

```
model2.fit(X_train, y_train)
```



```
▼ AdaBoostClassifier  
AdaBoostClassifier(random_state=42)
```

```
model3.fit(X_train, y_train)
```



```
▼ GradientBoostingClassifier  
GradientBoostingClassifier(n_estimators=50, random_state=42)
```

```
model4.fit(X_train, y_train)
```



```
▼ XGBClassifier  
XGBClassifier(base_score=None, booster=None, callbacks=None,  
               colsample_bylevel=None, colsample_bynode=None,  
               colsample_bytree=None, device=None, early_stopping_rounds=None,  
               enable_categorical=False, eval_metric='mlogloss',  
               feature_types=None, gamma=None, grow_policy=None,  
               importance_type=None, interaction_constraints=None,  
               learning_rate=None, max_bin=None, max_cat_threshold=None,  
               max_cat_to_onehot=None, max_delta_step=None, max_depth=None,  
               max_leaves=None, min_child_weight=None, missing=nan,  
               monotone_constraints=None, multi_strategy=None, n_estimators=50,  
               n_jobs=None, num_parallel_tree=None, objective='multi:softprob', ...)
```

✓ AdaBoostClassifier Evaluation

```
test_predictions = model2.predict(X_test)  
test_predictions
```



```
array([2, 0, 0, 1, 1, 2, 1, 0, 2, 1, 0, 2, 1, 1, 2, 1, 2, 2, 1, 2, 2, 2,  
       1, 2, 1, 1, 2, 1, 1, 1, 1, 0, 1, 2, 1, 0, 0, 1, 1, 0, 2, 2, 0, 0,  
       1, 2, 2, 2, 2, 1, 2, 1, 1, 2, 2, 2, 2, 1, 1, 1, 2, 2, 2, 1, 1, 0,  
       2, 1, 1, 0, 0, 0, 0, 2, 2, 0, 1, 2, 0, 2, 1, 1, 1, 1, 2, 1, 1, 0,  
       0, 2, 2, 0, 1, 2, 2, 1, 2, 1, 0, 2, 1, 0, 1, 1, 2, 2, 1, 2, 1, 2,  
       2, 1, 2, 2, 0, 1, 1, 0, 1, 0, 0, 1, 2, 1, 1, 0, 2, 2, 2, 2, 1, 1,  
       2, 2, 1, 2, 0, 2, 2, 1, 0, 2, 2, 2, 1, 2, 0, 1, 1, 2, 2, 2, 2, 2,  
       0, 2, 1, 1, 2, 0, 1, 1, 2, 2, 1, 2, 2, 1, 1, 2, 2, 1, 2, 1, 2, 0,  
       2, 1, 1, 1, 0, 2, 2, 0, 0, 1, 0, 2, 1, 2, 2, 1, 1, 2, 1, 1, 2, 0,  
       1, 1, 1, 2, 2])
```

```
accuracy2=accuracy_score(y_test, test_predictions)  
accuracy2
```



```
0.6945812807881774
```

```
precision2=precision_score(y_test, test_predictions, average='weighted')  
precision2
```



```
0.7036267138434321
```



```
recall2=recall_score(y_test, test_predictions, average='weighted')
recall2
```

```
0.6945812807881774
```

```
f12=f1_score(y_test, test_predictions, average='weighted')
f12
```

```
0.69643113617261
```

✓ Show accuracy matrix as a Dataframe

```
accuracy_matrix = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score"],
    "Score": [accuracy2, precision2, recall2, f12]
})
```

```
accuracy_matrix
```

```

   Metric  Score
0  Accuracy  0.694581
1  Precision  0.703627
2    Recall  0.694581
3  F1-Score  0.696431
```

✓ Compute confusion matrix


```
conf_matrix = confusion_matrix(y_test, test_predictions)
```

```
conf_matrix
```

```
array([[31,  3, 13],
       [ 2, 58, 20],
       [ 5, 19, 52]])
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True)
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
```

```
plt.show()
```



✓ GradientBoostingClassifier Evaluation

```
test_predictions = model3.predict(X_test)
test_predictions
```



```
array([2, 0, 0, 1, 1, 1, 1, 2, 2, 1, 0, 1, 1, 1, 2, 1, 0, 1, 1, 2, 2, 2,
       2, 0, 1, 2, 0, 1, 1, 1, 1, 0, 1, 2, 1, 0, 0, 2, 1, 0, 0, 2, 0, 0,
       1, 2, 0, 2, 1, 1, 1, 1, 1, 2, 2, 2, 0, 1, 1, 1, 1, 2, 1, 1, 0,
       2, 1, 2, 0, 0, 0, 2, 2, 2, 0, 1, 2, 0, 1, 1, 2, 1, 1, 2, 1, 1, 2,
       0, 2, 2, 2, 1, 0, 2, 1, 2, 1, 0, 2, 1, 0, 2, 1, 2, 2, 1, 2, 1, 2,
       2, 2, 1, 0, 2, 1, 1, 0, 2, 0, 2, 1, 1, 2, 1, 0, 2, 2, 1, 2, 1, 1,
       2, 2, 1, 1, 0, 2, 1, 2, 0, 2, 0, 0, 1, 1, 0, 1, 1, 2, 1, 2, 2, 2,
       0, 2, 1, 2, 1, 0, 1, 1, 1, 2, 1, 2, 0, 1, 1, 1, 0, 1, 2, 1, 1, 0,
       2, 1, 1, 1, 0, 1, 2, 0, 0, 1, 0, 1, 1, 1, 2, 2, 2, 2, 1, 1, 2, 0,
       1, 1, 2, 2, 0])
```

```
accuracy3=accuracy_score(y_test, test_predictions)
accuracy3
```



```
0.7487684729064039
```

```
precision3=precision_score(y_test, test_predictions, average='weighted')
precision3
```

➡ 0.7495390214602038

```
recall3=recall_score(y_test, test_predictions, average='weighted')
recall3
```

➡ 0.7487684729064039

```
f13=f1_score(y_test, test_predictions, average='weighted')
f13
```

➡ 0.7479471055637963

✓ Show accuracy matrix as a Dataframe

```
accuracy_matrix = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score"],
    "Score": [accuracy3, precision3, recall3, f13]
})
```

accuracy_matrix

➡

	Metric	Score
0	Accuracy	0.748768
1	Precision	0.749539
2	Recall	0.748768
3	F1-Score	0.747947

✓ Compute confusion matrix

```
conf_matrix = confusion_matrix(y_test, test_predictions)
```


conf_matrix

➡

```
array([[38,  4,  5],
       [ 1, 64, 15],
```

```
[ 6, 20, 50]])
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True)
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



✓ XGBClassifier Evaluation

```
test_predictions = model4.predict(X_test)
test_predictions
```



```
array([2, 0, 0, 1, 1, 2, 2, 2, 2, 1, 0, 2, 1, 0, 2, 2, 0, 1, 1, 2, 2, 2,
       0, 0, 1, 2, 0, 1, 1, 1, 2, 0, 1, 2, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0,
       1, 2, 0, 2, 1, 1, 2, 1, 1, 2, 2, 2, 0, 1, 1, 2, 1, 2, 2, 1, 1, 0,
       2, 0, 2, 0, 0, 0, 1, 2, 2, 0, 1, 2, 0, 2, 1, 2, 1, 1, 2, 1, 1, 2,
       0, 2, 2, 0, 1, 0, 2, 1, 2, 1, 0, 2, 1, 2, 2, 1, 2, 2, 1, 2, 1, 2,
       2, 2, 1, 0, 2, 1, 1, 0, 1, 0, 2, 1, 1, 2, 0, 0, 2, 2, 2, 2, 1, 1,
       2, 2, 1, 1, 0, 2, 1, 2, 0, 2, 0, 0, 1, 1, 0, 1, 1, 2, 1, 2, 1, 2,
       0, 2, 1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 2, 1, 1, 1, 0, 1, 2, 1, 2, 0,
```

```
2, 1, 1, 1, 0, 2, 1, 0, 0, 1, 0, 1, 1, 1, 2, 2, 2, 2, 1, 1, 2, 0,
1, 2, 2, 2, 0])
```

```
accuracy4=accuracy_score(y_test, test_predictions)
accuracy4
```

```
↔ 0.8423645320197044
```

```
precision4=precision_score(y_test, test_predictions, average='weighted')
precision4
```

```
↔ 0.8427400801169274
```

```
recall4=recall_score(y_test, test_predictions, average='weighted')
recall4
```

```
↔ 0.8423645320197044
```

```
f14=f1_score(y_test, test_predictions, average='weighted')
f14
```

```
↔ 0.8423566439324042
```

✓ Show accuracy matrix as a Dataframe

```
accuracy_matrix4 = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score"],
    "Score": [accuracy4, precision4, recall4, f14]
})
```

```
accuracy_matrix4
```

```
↔
```


	Metric	Score
0	Accuracy	0.842365
1	Precision	0.842740
2	Recall	0.842365
3	F1-Score	0.842357

✓ Compute confusion matrix

```
conf_matrix4 = confusion_matrix(y_test, test_predictions)
conf_matrix4
```

```
⇒ array([[41,  2,  4],
        [ 4, 66, 10],
        [ 3,  9, 64]])
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix4, annot=True)
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



- ✓ Choose base models among AdaBoostClassifier, GradientBoostingClassifier and XGBClassifier

```
models = {
    'Bagging': model1,
    'AdaBoost': model2,
    'GradientBoosting': model3,
```

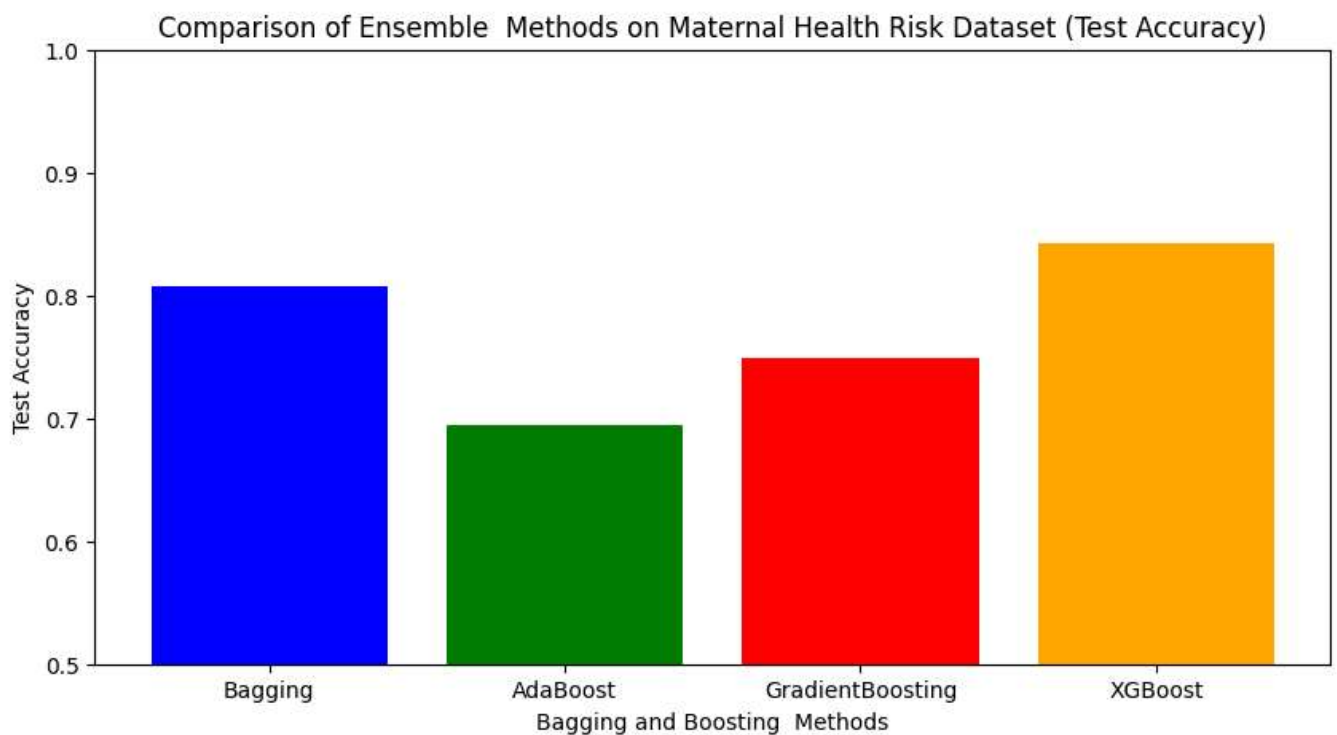
```
'XGBoost': model4
```

```
}
```

```
test_accuracies = {}
```

```
for name, model in models.items():  
    model.fit(X_train, y_train)  
    test_predictions = model.predict(X_test)  
    test_accuracies[name] = accuracy_score(y_test, test_predictions)
```

```
plt.figure(figsize=(10, 5))  
plt.bar(test_accuracies.keys(), test_accuracies.values(), color=['blue', 'green', 'red', 'orange'])  
plt.xlabel('Bagging and Boosting Methods')  
plt.ylabel('Test Accuracy')  
plt.title('Comparison of Ensemble Methods on Maternal Health Risk Dataset (Test Accuracy)')  
plt.ylim([0.5, 1.0]) # Adjusted for dataset accuracy range  
plt.show()
```



```
best_model_name = max(test_accuracies, key=test_accuracies.get)
```

```
best_model_name
```

 'XGBoost'

```
print("\nTest Accuracy:")
for model, acc in test accuracies.items():
    print(f"{model}: {acc:.4f}")
```



```
Test Accuracy:
Bagging: 0.8079
AdaBoost: 0.6946
GradientBoosting: 0.7488
XGBoost: 0.8424
```

```
print(f"\n✅ Best Model: **{best_model_name}**")
```



```
✅ Best Model: **XGBoost**
```

```
best_model = models[best_model_name]
```

```
best_model.fit(X_train, y_train)
best_test_predictions = best_model.predict(X_test)
```

Start coding or [generate](#) with AI.

```
best_test_accuracy = accuracy_score(y_test, best_test_predictions)
```

```
print(f"🎯 Best Model Test Accuracy: {best_test_accuracy:.4f}")
```



```
🎯 Best Model Test Accuracy: 0.8424
```

✓ 4) Stacking Implementation:

```
from sklearn.ensemble import StackingClassifier
```

```
from sklearn.tree import DecisionTreeClassifier
```

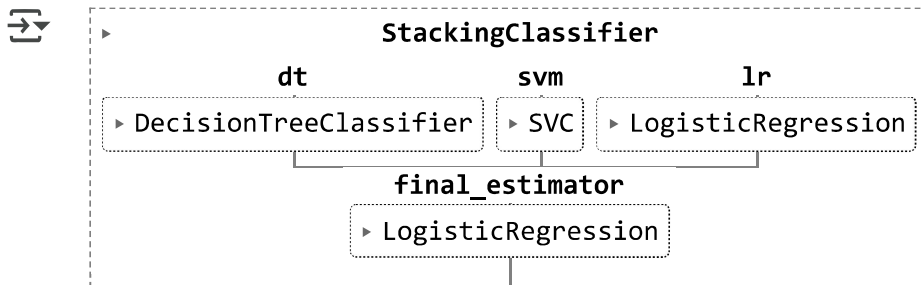
```
from sklearn.svm import SVC
```



```
from sklearn.linear_model import LogisticRegression
```

```
model5=StackingClassifier(  
    estimators=[('dt', DecisionTreeClassifier()), ('svm', SVC(probability=True)), ('lr',  
    final_estimator=LogisticRegression()  
    )
```

```
model5.fit(X_train, y_train)
```



```
test_predictions = model5.predict(X_test)  
test_predictions
```

```
array([[2, 0, 0, 1, 1, 2, 2, 2, 2, 1, 0, 2, 2, 0, 2, 2, 0, 1, 1, 2, 2, 2,  
       0, 0, 0, 2, 0, 1, 1, 1, 2, 0, 1, 2, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0,  
       1, 2, 0, 2, 1, 1, 2, 1, 1, 2, 2, 2, 0, 1, 1, 2, 1, 2, 2, 1, 1, 0,  
       2, 0, 2, 0, 0, 0, 2, 2, 2, 0, 0, 0, 0, 2, 1, 2, 1, 1, 2, 1, 1, 2,  
       0, 2, 2, 0, 1, 0, 2, 1, 2, 1, 0, 2, 1, 2, 2, 1, 2, 2, 1, 2, 1, 2,  
       2, 2, 1, 0, 2, 1, 1, 0, 1, 0, 2, 1, 1, 2, 0, 0, 2, 2, 2, 1, 1, 1,  
       2, 2, 1, 1, 0, 2, 1, 2, 0, 2, 0, 0, 1, 1, 0, 1, 1, 2, 1, 2, 1, 2,  
       0, 2, 1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 2, 2, 1, 1, 0, 1, 2, 1, 2, 0,  
       2, 1, 1, 1, 0, 2, 1, 2, 0, 1, 0, 1, 2, 1, 2, 2, 2, 2, 1, 2, 2, 0,  
       1, 2, 2, 2, 0]])
```

```
accuracy5=accuracy_score(y_test, test_predictions)  
accuracy5
```

```
0.8226600985221675
```

```
precision5=precision_score(y_test, test_predictions, average='weighted')  
precision5
```

```
0.8252832151323424
```

```
recall5=recall_score(y_test, test_predictions, average='weighted')  
recall5
```

```
0.8226600985221675
```

```
f15=f1_score(y_test, test_predictions, average='weighted')
f15
```

↗ 0.8225547363607623

✓ Show accuracy matrix as a Dataframe

```
accuracy_matrix = pd.DataFrame({
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score"],
    "Score": [accuracy5, precision5, recall5, f15]
})
```

accuracy_matrix

↗

	Metric	Score
0	Accuracy	0.822660
1	Precision	0.825283
2	Recall	0.822660
3	F1-Score	0.822555


✓ Compute confusion matrix

```
conf_matrix5 = confusion_matrix(y_test, test_predictions)
conf_matrix5
```

↗

```
array([[41,  1,  5],
       [ 5, 62, 13],
       [ 3,  9, 64]])
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True)
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



5) Compare the performance of Bagging, Boosting, and Stacking

```
performance={'Models':['Bagging','Boosting(XGB (best))','Stacking'],
             'Accuracy':[accuracy1,accuracy4,accuracy5],
             'Precision':[precision1,precision4,precision5],
             'Recall':[recall1,recall4,recall5],
             'F1-Score':[f11,f14,f15]}

}
```

```
pd.DataFrame(performance)
```



	Models	Accuracy	Precision	Recall	F1-Score
0	Bagging	0.807882	0.814527	0.807882	0.808429
1	Boosting(XGB (best))	0.842365	0.842740	0.842365	0.842357
2	Stacking	0.822660	0.825283	0.822660	0.822555

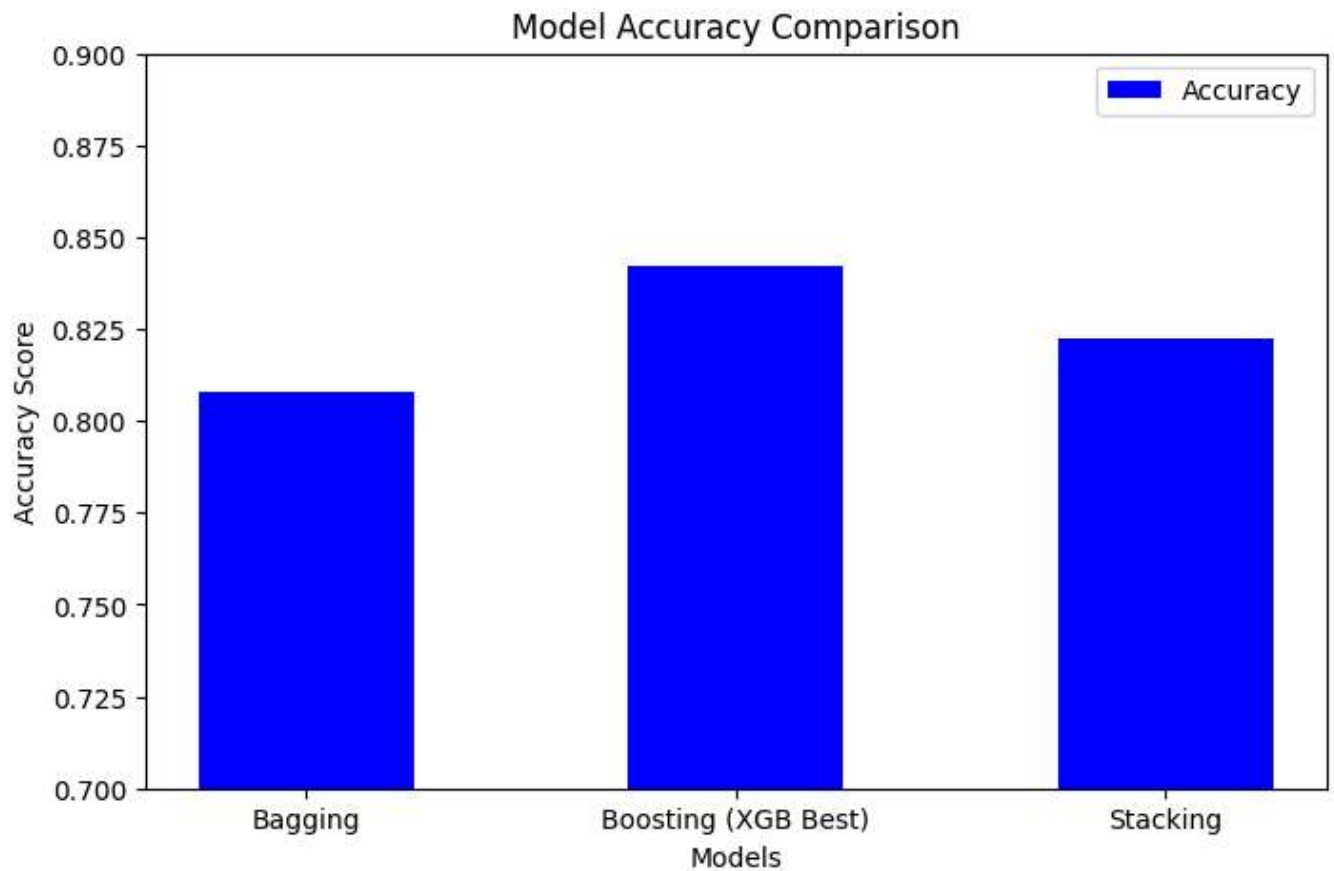
```
performance = {  
    'Models': ['Bagging', 'Boosting (XGB Best)', 'Stacking'],  
    'Accuracy': [accuracy1, accuracy4, accuracy5] # Only accuracy values  
}
```

```
models = performance['Models']  
accuracy_values = performance['Accuracy']
```

```
import numpy as np
```

```
x = np.arange(len(models))  
width = 0.5
```

```
fig, ax = plt.subplots(figsize=(8, 5))  
ax.bar(x, accuracy_values, width, color='blue', label='Accuracy')  
ax.set_xlabel('Models')  
ax.set_ylabel('Accuracy Score')  
ax.set_title('Model Accuracy Comparison')  
ax.set_xticks(x)  
ax.set_xticklabels(models)  
ax.legend()  
plt.ylim(0.7, 0.9)  
plt.show()
```




```
model_names = ['Bagging', 'Boosting (XGB Best)', 'Stacking']
```

```
conf_matrices = [conf_matrix1, conf_matrix4, conf_matrix5]
```

```
fig, axes = plt.subplots(1, 3, figsize=(15, 5))
for i, ax in enumerate(axes):
    sns.heatmap(conf_matrices[i], annot=True, fmt="d", cmap="Blues", cbar=False, ax=ax)
    ax.set_title(f'Confusion Matrix - {model_names[i]}')
    ax.set_xlabel('Predicted Label')
    ax.set_ylabel('True Label')

plt.tight_layout()

plt.show()
```



✓ 6) Select the Model (XGB Classifier) Analysis

```
accuracy_matrix4 = pd.DataFrame({  
    "Metric": ["Accuracy", "Precision", "Recall", "F1-Score"],  
    "Score": [accuracy4, precision4, recall4, f14]  
})
```


```
pd.DataFrame(accuracy_matrix4)
```



	Metric	Score
0	Accuracy	0.842365
1	Precision	0.842740
2	Recall	0.842365
3	F1-Score	0.842357

```
plt.figure(figsize=(6, 5))  
sns.heatmap(conf_matrix4, annot=True)  
plt.xlabel("Predicted Labels")
```

```
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



✓ XGB Classifier (Best) with Validation

```
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.2, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=
```

```
final_model=XGBClassifier(n_estimators=50, eval_metric='mlogloss', random_state=42)
```

```
final_model.fit(X_train, y_train)
```



```
▼ XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric='mlogloss',
               feature_types=None, gamma=None, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=None, max_bin=None, max_cat_threshold=None,
               max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
               max_leaves=None, min_child_weight=None, missing=nan,
               monotone_constraints=None, multi_strategy=None, n_estimators=50,
               n_jobs=None, num_parallel_tree=None, objective='multi:softprob', ...)
```

```
val_predictions = final_model.predict(X_val)
test_predictions = final_model.predict(X_test)
```

```
accuracy_score(y_test, test_predictions)
```




```
0.8529411764705882
```

```
conf_matrix = confusion_matrix(y_test, test_predictions)
conf_matrix
```



```
array([[23,  0,  3],
       [ 3, 32,  3],
       [ 1,  5, 32]])
```

```
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True)
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```

```
final_evaluation={'XGB':[accuracy4],  
                 'XGB(Validation)':[accuracy_score(y_test, test_predictions)]}
```

```
final_evaluation
```

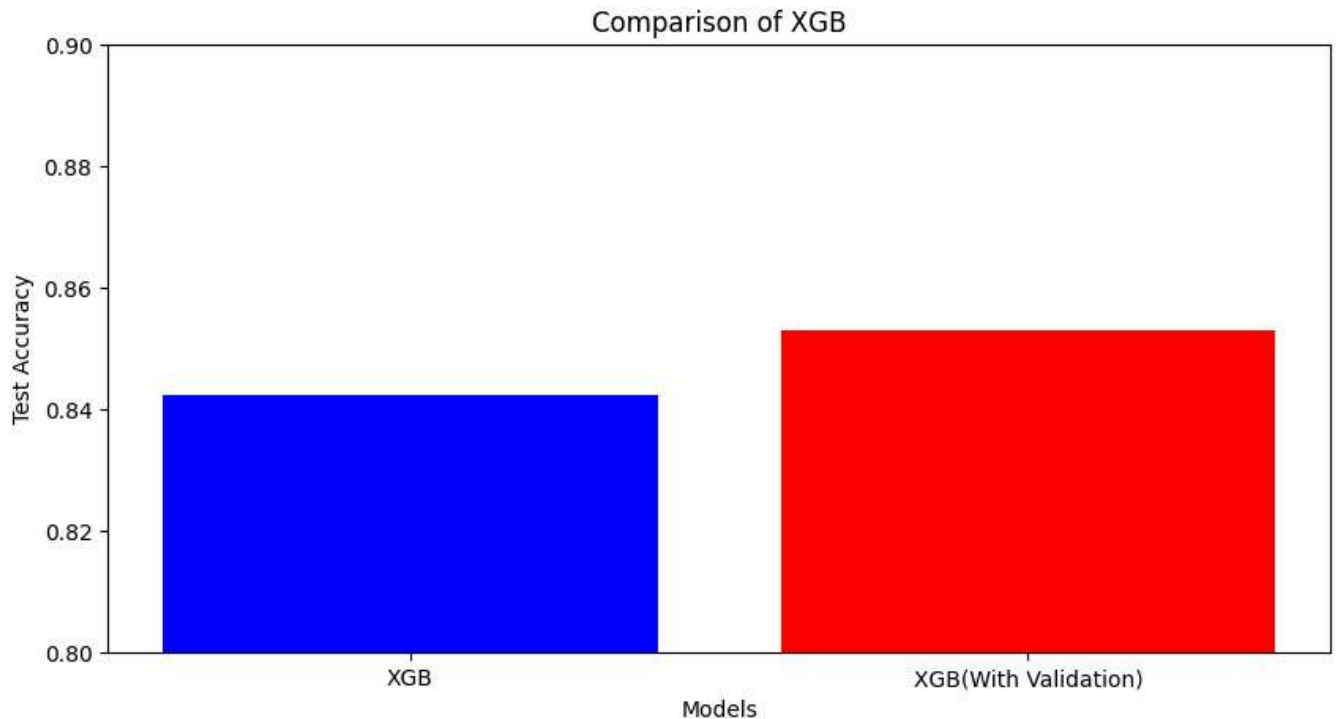


```
{'XGB': [0.8423645320197044], 'XGB(Validation)': [0.8529411764705882]}
```

```
final_evaluation = {  
    'XGB': [0.8423645320197044],  
    'XGB(With Validation)': [0.8529411764705882]  
}
```

```
models = list(final_evaluation.keys())  
accuracies = [val[0] for val in final_evaluation.values()]
```

```
plt.figure(figsize=(10, 5))
plt.bar(models, accuracies, color=['blue', 'red'])
plt.xlabel('Models')
plt.ylabel('Test Accuracy')
plt.title('Comparison of XGB')
plt.ylim([0.8, 0.9])
plt.show()
```



```
pd.DataFrame(final_evaluation )
```



	XGB	XGB(With Validation)
0	0.842365	0.852941


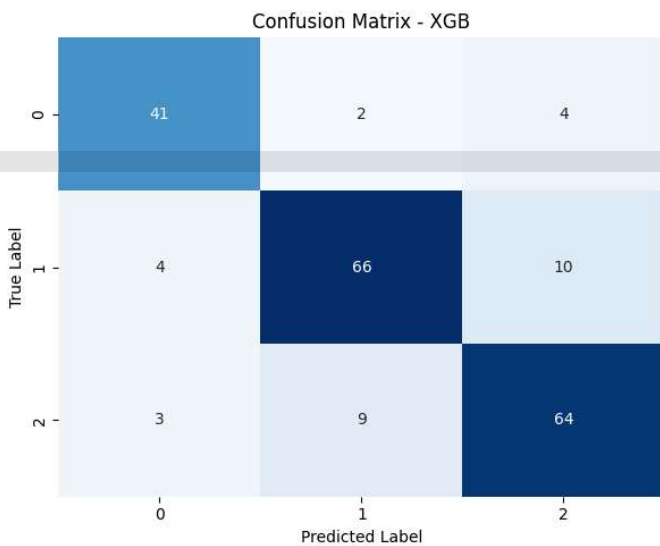
```
model_names = ['XGB', 'XGB (Validation)']
```

```
conf_matrices = [conf_matrix4, conf_matrix]
```

```
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
for i, ax in enumerate(axes):
    sns.heatmap(conf_matrices[i], annot=True, fmt="d", cmap="Blues", cbar=False, ax=ax)
    ax.set_title(f'Confusion Matrix - {model_names[i]}')
```

```
ax.set_xlabel('Predicted Label')
ax.set_ylabel('True Label')
```

```
plt.tight_layout()
```



✓ End of Assignment

Start coding or [generate](#) with AI.